



Diffusion Models

DL4DS – Spring 2025


Based on [Rocca, 2022, "Understanding Diffusion Probabilistic Models \(DPMs\)", Towards Data Science](#) and other references...

DS542 Gardos
Prince, [Understanding Deep Learning](#),
Other Content Cited

Outline

- Contextualizing Diffusion Models
- Theory behind diffusion models
- Architectures and Training
- Applications

Different Generative Models



[L. Weng, “What are Diffusion Models?”, blog post, 2021](#)

Given a probability distribution only described by some available samples, how can one generate a new sample?

Introduction

1D example:
we illustrate the
effet of G over
the entire
distribution

*Generative model
to be learned*
*Simple 1D gaussian
distribution we know
how to sample from*

$$G(\text{---}) = \text{---}$$

High dimension
example:
we illustrate the
effet of G over a
single sample

*Generative model
to be learned*
*High dimension data
point from simple
noise distribution*

$$G(\text{---}) = \text{---}$$

*Targeted complex 1D
distribution we don't know
how to sample from*



*High dimension data
point from complex
image distribution*

Generative models aims at learning a function that takes data from a simple distribution and transform it into data from a complex distribution.

Retrospective

- 2013: Kingma and Welling introduce Variational AutoEncoder.
- 2014: Goodfellow et al introduced Generative Adversarial Networks (GANs).
- 2015: Sohl-Dickstein “Deep Unsupervised Learning using Nonequilibrium Thermodynamics”
- 2020 (DDPM): J. Ho, A. Jain, and P. Abbeel, “Denoising Diffusion Probabilistic Models”
- 2022 (DDIM): J. Song, C. Meng, and S. Ermon, “Denoising Diffusion Implicit Models” – more efficient
- 2022 (Stable Diffusion): R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, “High-Resolution Image Synthesis with Latent Diffusion Models”
- 2023 (Diffusion Transformer): W. Peebles and S. Xie, “Scalable Diffusion Models with Transformers”


UDL Chapter 18 – Diffusion Models

- 18.2 Encoder (forward process)
 - Defines the noise diffusion process
- 18.3 Decoder model (reverse process)
 - Derives the denoising steps based on neural networks
- 18.4 Training
 - Derives a lower bound (ELBO) on the loss function
 - Initial diffusion loss function
- 18.5 Reparameterization of loss function
 - Shown to work better empirically
- 18.6 Implementaiton
 - Training and sampling algorithms

Commercial Diffusion Solutions

- [Dall-E 2](#) and 3 from OpenAI
- [Imagen](#) from Google
- [Make-A-Scene](#) from Meta
- [Imagen Video](#) from Google
- [Make-A-Video](#) from Meta
- [Stable Diffusion 3](#) from stability.ai
- [Model Version 6](#) from midjourney.com


Examples



Examples above have been generated by Meta Make-A-Scene model, that generates images from both a text prompt and a basic sketch for greater level of creative control.

Basic idea of Diffusion Probabilistic Models

- learn the *reverse process* of
- a well defined *stochastic forward process* that progressively destroys information, taking data from our complex target distribution and bringing them to a simple gaussian distribution.
- *reverse process* is then expected to take the path in the opposite direction, taking gaussian noise as an input and generating data from the distribution of interest.



Rocca, 2022

Outline

First:

- Stochastic Process
- Diffusion Process

Then intuition behind DPMs

Then some math basis

Then how trained in practice

Markov stochastic process


Stochastic Processes

- Discrete: $X_n, \quad \forall n \in \mathbb{N}$
- Continuous: $X_t, \quad \forall t \geq 0$


Realization of a random variable → sample

Realization of a stochastic process → sample path or trajectory

Different types of stochastic processes



Different types of stochastic processes



Rocca, 2022

14

Markov (Stochastic) Process

A *Markov process* is a *stochastic process* with no memory.

Future behavior only depends on the present, or

Present only depends on the previous sample.

$$P(X_{t_n} | X_{t_{n-1}}, \dots, X_{t_0}) = P(X_{t_n} | X_{t_{n-1}}) \quad \forall t_0 < t_1 < \dots < t_{n-1} < t_n$$

Diffusion Process

Any *diffusion process* can be described by a *stochastic differential equation* (SDE)

$$dX_t = a(X_t, t)dt + \sigma(X_t, t)dW_t$$

where:

$a(\cdot)$ is called the *drift coefficient*

$\sigma(\cdot)$ is called the *diffusion coefficient*

W is the *Wiener process*

Both a and σ are a function of the value and time

Diffusion Process

Any *diffusion process* can be described by a *stochastic differential equation* (SDE)

$$dX_t = a(X_t, t)dt + \sigma(X_t, t)dW_t$$

The equation is shown with two blue brackets. One bracket groups the term $a(X_t, t)dt$ and is labeled "Simple differential equation". Another bracket groups the term $\sigma(X_t, t)dW_t$ and is labeled "Stochastic part".

where:

$a(\cdot)$ is called the **drift coefficient**

$\sigma(\cdot)$ is called the **diffusion coefficient**


W is the **Wiener process**

Wiener Process (Brownian Motion)

Continuous time stochastic process

The Wiener process W_t is characterised by the following properties:^[2]

1. $W_0 = 0$ almost surely
2. W has independent increments: for every $t > 0$, the future increments $W_{t+u} - W_t$, $u \geq 0$, are independent of the past values W_s , $s < t$.
3. W has Gaussian increments: $W_{t+u} - W_t$ is normally distributed with mean 0 and variance u ,
→ $W_{t+u} - W_t \sim \mathcal{N}(0, u)$.
4. W has almost surely continuous paths: W_t is almost surely continuous in t .



Norbert Wiener




Norbert Wiener (November 26, 1894 – March 18, 1964) was an American computer scientist, mathematician and philosopher. He became a professor of mathematics at the Massachusetts Institute of Technology (MIT).

A child prodigy, Wiener later became *an early researcher in stochastic and mathematical noise processes*, contributing work relevant to electronic engineering, electronic communication, and control systems.

Wiener is considered the originator of cybernetics, the science of communication as it relates to living things and machines.

Heavily influenced John von Neumann, Claude Shannon, etc...

Wrote “The Machine Age” in 1949 anticipating robots, etc.



https://en.wikipedia.org/wiki/Norbert_Wiener

<https://www.nytimes.com/2013/05/21/science/mit-scholars-1949-essay-on-machine-age-is-found.html>

Discretizing

So

$$dW_t \approx W_{t+dt} - W_t \sim \mathcal{N}(0, dt)$$

Property of Weiner Process:
The std is equal to the time step.


Discretizing the SDE

$$X_{t+dt} - X_t \approx a(X_t, t)dt + \sigma(X_t, t)U \quad \text{where } U \sim \mathcal{N}(0, dt)$$

Which can also be rewritten

$$X_{t+dt} \approx X_t + \underbrace{a(X_t, t)dt}_{\text{Deterministic drift term}} + \underbrace{U'}_{\text{Normal RV with std proportional to diffusion term}} \quad \text{where } U' \sim \mathcal{N}(0, \underbrace{\sigma(X_t, t)dt}_{\text{diffusion term}})$$

Diffusion process samples



Reversed time process

If X_t is a diffusion process such that

$$dX_t = a(X_t, t)dt + \sigma(t)dW_t$$

then the reversed-time process, $\bar{X}_t = X_{T-t}$ is also a diffusion process

$$\begin{aligned} d\bar{X}_t &= [a(\bar{X}_t, t) - \sigma^2(t)\nabla_{X_t} \log p(X_t)]dt + \sigma(t)dW_t \\ &= \bar{a}(\bar{X}_t, t)dt + \sigma(t)dW_t \end{aligned}$$


where $\nabla_{X_t} \log p(X_t)$ is called the *score function* and $p(X_t)$ is the *marginal probability* of X_t

Intuition behind diffusion processes

Progressively destroys relevant information

E.g. with shrinking ($|a| < 1$) *drift coefficient* and non-zero *diffusion coefficient* will turn complex distribution into isotropic gaussian


$$X_t = \sqrt{1-p}X_{t-1} + \sqrt{p}u_t \quad \text{where} \quad u_t \sim \mathcal{N}(0, 1) \quad \text{and} \quad p = 0.01$$



Intuition behind diffusion processes

For the diffusion process

$$X_t = \sqrt{1-p}X_{t-1} + \sqrt{p}u_t \quad \text{where} \quad u_t \sim \mathcal{N}(0, 1) \quad \text{and} \quad p = 0.01$$




After a given number of steps T we can write

$$\begin{aligned} X_T &= \sqrt{1-p}X_{T-1} + \sqrt{p}u_T \\ &= (\sqrt{1-p})^2 X_{T-2} + \sqrt{1-p}\sqrt{p}u_{T-1} + \sqrt{p}u_T \\ &= \dots \\ &= (\sqrt{1-p})^T X_0 + \sum_{i=0}^{T-1} \underbrace{\sqrt{p}(\sqrt{1-p})^i}_{\text{This is a sum of independent gaussians}} u_{T-i} \end{aligned}$$

This is a sum of independent gaussians,
so can express as single gaussian with
variance the sum of the variances.

Intuition behind diffusion processes



Towards Gaussian

$$(\sqrt{1-p})^T \xrightarrow{T \rightarrow \infty} 0 \quad \text{and}$$

After a given number of steps T we can write

$$\begin{aligned} X_T &= \sqrt{1-p}X_{T-1} + \sqrt{p}u_T \\ &= (\sqrt{1-p})^2X_{T-2} + \sqrt{1-p}\sqrt{p}u_{T-1} + \sqrt{p}u_T \\ &= \dots \\ &= (\sqrt{1-p})^T X_0 + \sum_{i=0}^{T-1} \sqrt{p}(\sqrt{1-p})^i u_{T-i} \end{aligned}$$

For number of steps T large enough, we have


$$\sum_{i=0}^{T-1} \underbrace{\left(\sqrt{p}(\sqrt{1-p})^i\right)^2}_{\text{Geometric series}} = p \frac{1 - (1-p)^T}{1 - (1-p)} \xrightarrow{T \rightarrow \infty} 1$$

Variance of the gaussian

So for any starting point, we tend to a normal gaussian.

Same idea but for images

But in $H \times W \times C$ dimensions, e.g. $100 \times 100 \times 3$ for 100×100 resolution RGB images



$$X_1 = \sqrt{1-p} X_0 + \sqrt{p} u_1 \sim \mathcal{N}(0, I)$$




The equation shows the generation of X_1 from X_0 . It is equal to the product of $\sqrt{1-p}$ and X_0 , plus the product of \sqrt{p} and u_1 , where u_1 follows a normal distribution $\mathcal{N}(0, I)$.

Why use diffusion?

Answer:


Gives us a progressive and structured way to go from a complex distribution to an isotropic gaussian noise
that will enable the learning of the reverse process

Intuition behind learning the reverse process




Reversing process in one step is extremely difficult

Doing it in steps gives us some clues




One step versus multi-step



Isn't it equivalent to learn a single big function G^* instead of learning T smaller functions G_i ?


If we set $G^* = G_1 \circ G_2 \circ \dots \circ G_{T-1} \circ G_T$ both tasks can look pretty similar at first sight.

Advantage of multi-step reverse process



1. Don't have to learn a unique transform G_i for each step, but rather a single transform that is a function of the index step. Drastically reduces size of the model.
2. Gradient descent is much more difficult in one step and can exploit coarse to fine adjustments in multiple steps.,

Iterative versus one step



G* (to be learned)

G* can't rely on the same nice iterative structure than G, meaning that this unrolled version supposed to be equivalent to $\mathbf{G}_1 \circ \mathbf{G}_2 \circ \dots \circ \mathbf{G}_{T-1} \circ \mathbf{G}_T$ will have more parameters and will be harder to train

Diffusion Model from Scratch

- HuggingFace Notebook
- <https://github.com/huggingface/diffusion-models-class>

[02_diffusion_models_from_scratch.ipynb](#)

Similarities and differences to VAEs


Similarities:

- An encoder transforms a complex distribution into a simple distribution in a structured way to learn a decoder that produces a similar sample

Differences:

- DPM is multi-step, versus one step for VAE
- DPM encoder is fixed and does not get trained
- DPM will be trained based on the structure of the diffusion process
- DPM latent space is exactly same as input, as opposed to VAE which reduces dimensionality

Dall-E 3



CAUTION:
MATHEMATICS
AHEAD

Mathematics of Diffusion Models

Assume the forward and reverse process operate in T steps.

Both forward and reverse process are discrete so becomes a *Markov chain with gaussian transition probability*.

Diffusion Process

Any *diffusion process* can be described by a *stochastic differential equation* (SDE)

$$dX_t = a(X_t, t)dt + \sigma(X_t, t)dW_t$$

where:

$a(\cdot)$ is called the *drift coefficient*

$\sigma(\cdot)$ is called the *diffusion coefficient*

W is the *Wiener process*

Both a and σ are a function of the value and time

16

$$\mathcal{N}(\mu, \sigma^2)$$

Mathematics of Diffusion Models

Denote x_0 as a sample from a distribution $q(x_0)$.

Forward process: gaussian transition probability

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{(1 - \beta_t)} x_{t-1}, \beta_t I) \quad \text{where } t \in \mathbb{N}$$

and where β_t indicates trade-off between info to be kept from previous step and new noise added.

$$\mathcal{N}(\mu, \sigma^2)$$

Mathematics of Diffusion Models

Denote x_0 as a sample from a distribution $q(x_0)$.

Forward process: gaussian transition probability

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{(1 - \beta_t)} x_{t-1}, \beta_t I) \quad \text{where } t \in \mathbb{N}$$

and where β_t indicates trade-off between info to be kept from previous step and new noise added.

We can equivalently write

$$x_t = \sqrt{(1 - \beta_t)} x_{t-1} + \sqrt{\beta_t} \epsilon_t \quad \epsilon_t \sim \mathcal{N}(0, I)$$

Discretized diffusion process

$$\mathcal{N}(\mu, \sigma^2)$$

Mathematics of Diffusion Models

Through recurrence, we can represent any step in the chain as directly represented from x_0 :

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t)I)$$

where

$$\alpha_t = (1 - \beta_t) \quad \text{and} \quad \bar{\alpha}_t = \prod_{i=1}^t \alpha_i = \prod_{i=1}^t (1 - \beta_i)$$

and from the Markov property, the entire forward trajectory is

$$q(x_{0:T}) = q(x_0) \prod_{t=1}^T q(x_t|x_{t-1})$$

The reverse process

With the assumption on the drift and diffusion coefficients, the reverse of the diffusion process takes the same form.

Reverse gaussian transition probability

$$q(x_{t-1} | x_t)$$

can then be approximated by

$$p_\theta(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

where μ_θ and Σ_θ are two functions parameterized by θ and learned.

The reverse process

Using the Markov property, the probability of a given backward trajectory can be approximated by

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t)$$

where $p(x_T)$ is an isotropic gaussian distribution that does not depend on θ

$$p(x_T) = \mathcal{N}(x_T; 0, I)$$


FIXED FORWARD PROCESS

Initial distribution

$$q(x_0)$$

Gaussian transition kernel

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$



Approximation of

$$q(x_{t-1}|x_t)$$

Gaussian transition kernel with parameters to be learned

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

Initial distribution

$$p(x_T) = \mathcal{N}(x_T; 0, I)$$

LEARNED BACKWARD PROCESS

Questions

How do we learn the parameters θ for μ_θ and Σ_θ ?

What is the loss to be optimized?

- We hope that $p_\theta(x_0)$, the distribution of the last step of the reverse process, will be close to $q(x_0)$

Optimization Objective

$$\begin{aligned}\mu_\theta^*, \Sigma_\theta^* &= \arg \min_{\mu_\theta, \Sigma_\theta} (D_{KL}(q(x_0) || p_\theta(x_0))) \\ &= \arg \min_{\mu_\theta, \Sigma_\theta} \left(- \int q(x_0) \log \left(\frac{p_\theta(x_0)}{q(x_0)} \right) dx_0 \right) \\ &= \arg \min_{\mu_\theta, \Sigma_\theta} \left(- \int q(x_0) \log(p_\theta(x_0)) dx_0 \right)\end{aligned}$$

Skipping a lot more math

- Expand p -theta as marginalization integral
- Use Jensen's inequality to define a slightly simpler upper bound to the loss
- Some manipulations with Bayes' Theorem
- Properties of KL divergence of two gaussian distributions
- An additional simplification suggested by [Ho et al 2020]

Diffusion models in practice

We have the forward process

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1-\beta_t}x_{t-1}, \beta_t I) \quad q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1-\bar{\alpha}_t)I)$$

and our reverse process

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

and we want to train to minimize this simplified upper bound

$$\mathbb{E}_{x_0, t, \epsilon} (||\epsilon - \epsilon_\theta(x_t, t)||^2) = \mathbb{E}_{x_0, t, \epsilon} (||\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1-\bar{\alpha}_t}\epsilon, t)||^2)$$



Dall-E 3 47


Training Process

Algorithm 1 Training

```

1: repeat
2:    $x_0 \sim q(x_0)$                                  $\triangleright$  Sample random initial data
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$              $\triangleright$  Sample random step
4:    $\epsilon \sim \mathcal{N}(0, I)$                        $\triangleright$  Sample random noise
5:    $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$      $\triangleright$  Rand. step of rand. trajectory
6:   Take gradient descent step on  $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(x_t, t)\|^2$      $\triangleright$  Optimisation
7: until converged

```



To sample/generate




Illustration of the sampling process of a denoising diffusion probabilistic model.

To sample/generate

Algorithm 2 Sampling


```
1:  $x_T \sim \mathcal{N}(0, I)$                                      ▷Initial isotropic gaussian noise sampling  
2: for  $t = T, \dots, 1$  do  
3:    $z \sim \mathcal{N}(0, I)$  if  $t > 1$  else  $z = 0$           ▷Sample random noise (if not last step)  
4:    $\tilde{\epsilon} = \epsilon_\theta(x_t, t)$                       ▷Estimated noise in current noisy data  
5:    $\tilde{x}_0 = \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{1 - \bar{\alpha}_t}\tilde{\epsilon})$     ▷Estimated  $x_0$  from estimated noise  
6:    $\tilde{\mu} = \mu_t(x_t, \tilde{x}_0) \left(= \frac{1}{\sqrt{\bar{\alpha}_t}} \left( x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) \right)$     ▷Mean for previous step sampling  
7:    $x_{t-1} = \tilde{\mu} + \sigma_t z$                          ▷Previous step sampling  
8: end for  
9: return  $x_0$ 

---



```

U-Net (2016) as the Model

Output is same size as the input.



U-Net for reverse diffusion



Conditional generation

Classifier Guidance

- Modify the denoising update from z_t to z_{t-1} to incorporate class information, c .

Guidance from text

- Condition on a sentence embedding computed from a language model

Conditional generation using classifier guidance



Figure 18.12 Conditional generation using classifier guidance. Image samples conditioned on different ImageNet classes. The same model produces high quality samples of highly varied image classes. Adapted from Dhariwal & Nichol (2021).

Cascaded conditional generation based on a text prompt




Figure 18.11 Cascaded conditional generation based on a text prompt. a) A diffusion model consisting of a series of U-Nets is used to generate a 64×64 image. b) This generation is conditioned on a sentence embedding computed by a language model. c) A higher resolution 256×256 image is generated and conditioned on the smaller image *and* the text encoding. d) This is repeated to create a 1024×1024 image. e) Final image sequence. Adapted from Saharia et al. (2022b).

Conditional generation using text prompts





Figure 18.13 Conditional generation using text prompts. Synthesized images from a cascaded generation framework, conditioned on a text prompt encoded by a large language model. The stochastic model can produce many different images compatible with the prompt. The model can count objects and incorporate text into images. Adapted from Saharia et al. (2022b).

Stable Diffusion – Latent Diffusion Models

Project the original data to a smaller latent space using a conventional autoencoder and then run the diffusion process in the smaller space.



For practice

- Notebook 18.1 – Diffusion Encoder in 1D



 Open in Colab

- Notebook 18.2 – Training Decoder



 Open in Colab

- Notebook 18.3 – 1D Reparameterized Model (more robust)



 Open in Colab

- Notebook 18.4 – Families of Diffusion Models (DDIM)




 Open in Colab

Resources

- UDL, Chapter 18, Diffusion Models
- [Rocca, "Understanding Diffusion Probabilistic Models \(DPMs\)", Medium.com, 2022](#)
- HuggingFace [Diffusion Models Class](#)
- CVPR 2023 Tutorial: Denoising Diffusion Models: A Generative Learning Big Bang, <https://cvpr.thecvf.com/virtual/2023/tutorial/18546>

For discussion sections:

- Recap the theory
- Show example code with MNIST
- Review Diffusion Transformer (time permitting)



[Link](#)